

# Stellar Consensus Protocol

Saravanan Vijayakumaran  
sarva@ee.iitb.ac.in

Department of Electrical Engineering  
Indian Institute of Technology Bombay

September 25, 2018

# Lecture Plan

- Consensus Protocol Terminology
- Related Protocols for Context
  - Paxos
  - PBFT
- Federated Byzantine Agreement Model
- Federated Voting
- Stellar Consensus Protocol (in brief)

# Consensus Protocol Terminology

- **Agents:** Parties interested in achieving consensus
- Each agent has an input
- Agents use protocol to agree on one of the inputs
- Each agent decides on a chosen value
- Agent failure modes
  - Stopping failure
  - Byzantine failure
- **Safety**
  - **Agreement:** No two non-faulty agents decide on different values
  - **Validity:** If all non-faulty agents have the same input  $v$ , then  $v$  is the only possible decision value
- **Liveness**
  - **Termination:** All non-faulty agents eventually decide
- Asynchronous network model
  - Messages may be delayed, duplicated, lost, reordered
  - No corrupted messages

Paxos

# Paxos

- Consensus protocol for non-Byzantine agents and asynchronous network
- Proposed by Leslie Lamport in 1989
- Number of agents is known
- Agents act as proposers, acceptors, or learners (multiple roles allowed)
- Proposers propose values
- Acceptors accept a value if requested by a proposer
- Once a majority of acceptors has accepted a value, consensus has been achieved
- Learners are interested in learning about consensus values
- Challenges
  - Messages indicating acceptance may be lost
  - Consensus may be achieved without proposers finding out
  - Multiple proposers may be simultaneously proposing values

# Paxos Protocol Phase 1

- Proposal made by proposers have a proposal number  $n$  from a totally ordered set
- Phase 1
  - Proposer sends a **prepare** request with number  $n$  to all acceptors
  - If acceptor receives a prepare request with number higher than any other previous prepare request, then
    1. it promises to not accept any more proposals with number less than  $n$  and
    2. returns highest-numbered proposal value (if any) it has accepted
- Example

Prop. No.	Value	Agent 1	Agent 2	Agent 3
1	7	7	$\langle \rangle$	$\langle \rangle$
2	8	8	$\langle \rangle$	$\langle \rangle$
3	9	$\langle \rangle$	$\langle \rangle$	9

For proposal 4, highest-numbered proposal accepted among all responses is used

# Paxos Protocol Phase 2

- Phase 2

- If proposer receives a response to its prepare request from a majority of acceptors, then it **either**
  - sends an **accept** request to each these acceptors with value  $v$  which is the highest-numbered proposal among the responses or
  - sends an **accept** request with any value if responses reported no proposals.
- If acceptor receives an accept request for a proposal number  $n$ , it accepts the proposal unless it has already responded to a prepare request having number greater than  $n$ .

- Example 1

Prop. No.	Value	Agent 1	Agent 2	Agent 3
1	7	7	$\langle \rangle$	$\langle \rangle$
2	8	8	$\langle \rangle$	$\langle \rangle$
3	9	$\langle \rangle$	$\langle \rangle$	9

- For proposal 4, proposer can send accept request with
  - 8 if only agents 1 and 2 respond
  - 9 if only agents 2 and 3 respond

# Paxos Protocol Phase 2

- Phase 2

- If proposer receives a response to its prepare request from a majority of acceptors, then it **either**
  - sends an **accept** request to each these acceptors with values  $v$  which is the highest-numbered proposal among the responses or
  - sends an **accept** request with any value if responses reported no proposals.
- If acceptor receives an accept request for a proposal number  $n$ , it accepts the proposal unless it has already responded to a prepare request having number greater than  $n$ .

- Example 2

Prop. No.	Value	Agent 1	Agent 2	Agent 3
1	8	8	$\langle \rangle$	$\langle \rangle$
2	9	9	$\langle \rangle$	9
3	9	$\langle \rangle$	$\langle \rangle$	9

- For proposal 4, proposer can send accept request with only value 9



# Paxos Protocol

- Phase 1

- Proposer sends a **prepare** request with number  $n$  to all acceptors
- If acceptor receives a prepare request with number higher than any other previous prepare request, then
  1. it promises to not accept any more proposals with number less than  $n$  and
  2. returns highest-numbered proposal value (if any) it has accepted

- Phase 2

- If proposer receives a response to its prepare request from a majority of acceptors, then it **either**
  - sends an **accept** request to each these acceptors with values  $v$  which is the highest-numbered proposal among the responses or
  - sends an **accept** request with any value if responses reported no proposals.
- If acceptor receives an accept request for a proposal number  $n$ , it accepts the proposal unless it has already responded to a prepare request having number greater than  $n$ .
- Learners need messages from a majority of acceptors to find out about consensus value

# Proposer Selection

- Lamport describes a method using timeouts
  - Each agent broadcasts its ID and the one with the highest ID is the proposer
- Presence of multiple proposers cannot violate safety but can affect liveness
  - Proposer  $p$  completes phase 1 for proposal number  $n_1$
  - Proposer  $q$  completes phase 1 for proposal number  $n_2 > n_1$
  - Proposer  $p$ 's phase 2 messages are ignored
  - Proposer  $p$  completes phase 1 for new proposal with number  $n_3 > n_2$
  - Proposer  $q$ 's phase 2 messages are ignored
  - And so on
- **FLP Impossibility Theorem:** No deterministic consensus algorithm can guarantee all three of safety, liveness, and fault-tolerance in an asynchronous system.

# Practical Byzantine Fault Tolerance

# PBFT

- Proposed in 1999 as an algorithm for state machine replication
  - Each agent is a replica of a state machine
  - Replicas need to achieve consensus on state transitions
- Assumes Byzantine agent failures and weak synchrony
  - Messages may be delayed, duplicated, lost, reordered
  - Delays do not grow faster than  $t$  indefinitely
- Guarantees safety and liveness if at most  $\lfloor \frac{n-1}{3} \rfloor$  out of  $n$  replicas are faulty
  - For  $f$  faulty replicas,  $3f + 1$  is the minimum number of replicas required
- Let  $\mathcal{R}$  be the set of replicas with cardinality  $3f + 1$
- Each replica is identified using an integer in  $0, 1, \dots, |\mathcal{R}| - 1$
- The algorithm moves through a sequence of **views**
- Views are numbered sequentially
- In view  $v$ , replica with identity  $v \bmod |\mathcal{R}|$  is the **primary** and the remaining replicas are **backups**

# PBFT Algorithm

- Rough outline
  1. A client sends a request to the primary to invoke a state machine operation
  2. Primary multicasts the request to the backups
  3. Replicas execute the request and send a reply to the client
  4. The client waits for  $f + 1$  replies from different replicas with same result
- Three phases in case of non-faulty primary
  - Pre-prepare
  - Prepare
  - Commit
- Pre-prepare phase
  - Primary in view  $v$  receives client request  $m$
  - Primary assigns a sequence number  $n$  to  $m$
  - Primary multicasts PRE-PREPARE message with  $m, v, n$  to all backups
  - Backup accepts PRE-PREPARE message if
    - it is in view  $v$  and
    - it has not accepted a PRE-PREPARE message for view  $v$  and sequence number  $n$  with different request

# PBFT Prepare Phase

- Prepare
  - If backup  $i$  accepts the PRE-PREPARE message, it enters the prepare phase
  - Multicasts PREPARE message with  $v, n, m, i$  to all other replicas
  - Adds both PRE-PREPARE and PREPARE messages to its log
- Define predicate **prepared**( $m, v, n, i$ ) to be true if and only if replica  $i$  has inserted in its log
  1. a PRE-PREPARE message with  $m, v, n$ , and
  2. at least  $2f$  PREPARE messages for  $m, v, n$ .
- Guarantees that non-faulty replicas agree on total order of requests in a view
  - **Invariant:** If **prepared**( $m, v, n, i$ ) is true, then **prepared**( $m', v, n, j$ ) is false for any non-faulty replica  $j$  where  $m' \neq m$
  - **prepared**( $m, v, n, i$ ) true  $\implies$  at least  $f + 1$  non-faulty replicas have sent PREPARE or PRE-PREPARE messages for  $m, v, n$
  - **prepared**( $m', v, n, j$ ) true  $\implies 2f + 1$  replicas have sent PREPARE or PRE-PREPARE messages for  $m', v, n$  to  $j$
  - At least one non-faulty replica has sent conflicting PREPAREs or PRE-PREPAREs  $\implies$  contradiction

# PBFT Commit Phase

- Commit
  - When **prepared**( $m, v, n, i$ ) becomes true, replica  $i$  multicasts a COMMIT message for  $m, v, n, i$
  - Replicas accept COMMIT messages which match their view and insert them into their logs
  - Replica  $i$  executes the operation requested by  $m$  when **committed-local**( $m, v, n, i$ ) becomes true and all requests with lower sequence number have been executed
- **committed-local**( $m, v, n, i$ ) is true if and only if
  1. **prepared**( $m, v, n, i$ ) is true and
  2. replica  $i$  has accepted  $2f + 1$  COMMITs (including its own) for  $m, v, n$
- **committed**( $m, v, n$ ) is true if and only if **prepared**( $m, v, n, j$ ) is true for all  $j$  in some set of  $f + 1$  non-faulty replicas
- **Invariant:** If **committed-local**( $m, v, n, i$ ) is true for some non-faulty  $i$ , then **committed**( $m, v, n$ ) is true
- At non-faulty replicas  $i$  and  $j$ , **committed-local**( $m, v, n, i$ ) and **committed-local**( $m', v, n, j$ ) cannot both be true for  $m \neq m'$

# PBFT View Change


- View changes are required when primary replica fails
- View-change algorithm
  1. If client does not receive replies before a timeout, it broadcasts the request to all replicas
  2. If request has already been processed, the replicas resend the reply to client
  3. If request was not received from primary, a backup starts a timer upon receiving the client's request
  4. If the timer expires while waiting for same request from primary, the backup multicasts a view-change message to all replicas
  5. When primary of view  $v + 1$  receives  $2f$  view-change messages, it multicasts a new-view message and enters view  $v + 1$



# Federated Byzantine Agreement

# Federated Byzantine Agreement

- **Definition:** An **federated Byzantine agreement system (FBAS)** is a pair  $\langle \mathbf{V}, \mathbf{Q} \rangle$  comprising of a set of nodes  $\mathbf{V}$  and a quorum function  $\mathbf{Q} : \mathbf{V} \mapsto 2^{2^{\mathbf{V}}} \setminus \{\emptyset\}$  specifying one or more quorum slices for each node, where a node belongs to all of its own quorum slices, i.e.  $\forall v \in \mathbf{V}, \forall q \in \mathbf{Q}(v), v \in q$ .
- Example

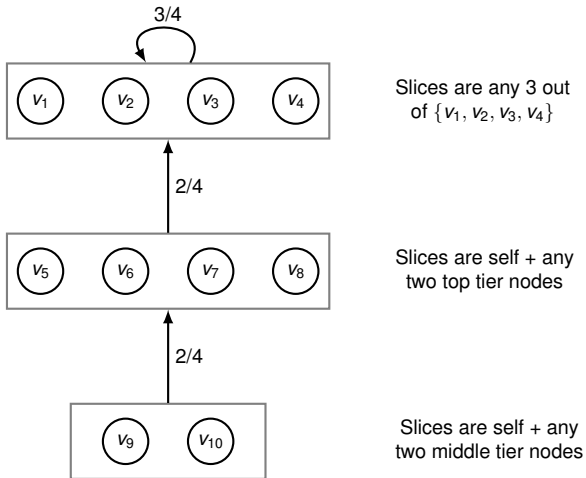


$$\mathbf{Q}(v_1) = \{\{v_1, v_2, v_3\}\}$$

$$\mathbf{Q}(v_2) = \mathbf{Q}(v_3) = \mathbf{Q}(v_4) = \{\{v_2, v_3, v_4\}\}$$

- **Definition:** A set of nodes  $\mathbf{U} \subseteq \mathbf{V}$  in FBAS  $\langle \mathbf{V}, \mathbf{Q} \rangle$  is a **quorum** iff  $\mathbf{U} \neq \emptyset$  and  $\mathbf{U}$  contains a slice for each member, i.e.  $\forall v \in \mathbf{U}, \exists q \in \mathbf{Q}(v)$  such that  $q \subseteq \mathbf{U}$ .
- A quorum of nodes is sufficient to reach agreement

# Tiered FBAS Example



Possible quorums?

# Safety and Liveness


- FBA systems attempt consensus in a slot
- A node applies update  $x$  in slot  $i$  when
  1. it has applied updates in all previous slots and
  2. it believes all non-faulty nodes will eventually agree on  $x$  for slot  $i$ .

The node is said to have **externalized**  $x$  in slot  $i$ .


- **Definition:** A set of nodes in an FBAS enjoy **safety** if no two of them ever externalize different values for the same slot
- Well-behaved nodes = obey protocol
- Ill-behaved nodes = Byzantine failures
- Well-behaved nodes can also fail (be blocked or diverge)
- **Definition:** A node in an FBAS enjoys **liveness** if it can externalize new values without the participation of any failed nodes
- Given a specific  $\langle \mathbf{V}, \mathbf{Q} \rangle$  and a ill-behaved subset of  $\mathbf{V}$ , what is the best any FBA protocol can do?

# Quorum Intersection

- **Definition:** An FBAS enjoys **quorum intersection** if and only if any two quorums share a node.
- No protocol can guarantee safety in absence of quorum intersection
- Example of quorum non-intersection




$$\begin{aligned} Q(v_1) &= Q(v_2) = Q(v_3) \\ &= \{\{v_1, v_2, v_3\}\} \end{aligned}$$



$$\begin{aligned} Q(v_4) &= Q(v_5) = Q(v_6) \\ &= \{\{v_4, v_5, v_6\}\} \end{aligned}$$

- $\{v_1, v_2, v_3\}$  and  $\{v_4, v_5, v_6\}$  are two disjoint quorums; can approve contradictory statements

# Quorum Intersection at Ill-Behaved Nodes



$$\begin{aligned}\mathbf{Q}(v_1) &= \mathbf{Q}(v_2) = \mathbf{Q}(v_3) \\ &= \{\{v_1, v_2, v_3, v_7\}\}\end{aligned}$$

$$\begin{aligned}\mathbf{Q}(v_4) &= \mathbf{Q}(v_5) = \mathbf{Q}(v_6) \\ &= \{\{v_4, v_5, v_6, v_7\}\}\end{aligned}$$

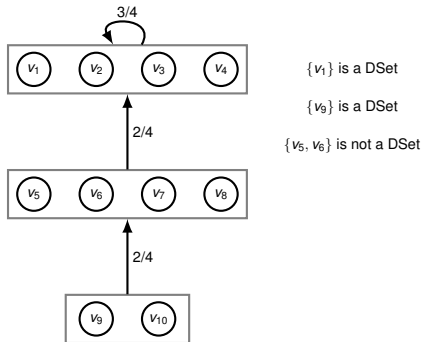
- If  $v_7$  is ill-behaved, the quorums are effectively disjoint
- **Necessary property for safety:** Well-behaved nodes enjoy quorum intersection after deleting ill-behaved nodes
- **Definition:** If  $\langle \mathbf{V}, \mathbf{Q} \rangle$  is an FBAS and  $B \subseteq \mathbf{V}$  is a set of nodes, to **delete**  $B$  is to compute the modified FBAS  $\langle \mathbf{V}, \mathbf{Q} \rangle^B = \langle \mathbf{V} \setminus B, \mathbf{Q}^B \rangle$  where  $\mathbf{Q}^B = \{q \setminus B \mid q \in \mathbf{Q}(v)\}$

# Dispensible Sets

- Safety and liveness of nodes outside a DSet can be guaranteed irrespective of the behaviour of nodes in the DSet
- **Definition:** Let  $\langle \mathbf{V}, \mathbf{Q} \rangle$  be an FBAS and  $B \subseteq \mathbf{V}$  be a set of nodes. We say  $B$  is a dispensible set or DSet if and only if
  1.  $\langle \mathbf{V}, \mathbf{Q} \rangle^B$  enjoys quorum intersection, and
  2. either  $\mathbf{V} \setminus B$  is a quorum in  $\langle \mathbf{V}, \mathbf{Q} \rangle$  or  $B = \mathbf{V}$ .

Condition 1 = *quorum intersection despite B*

Condition 2 = *quorum availability despite B*



## Intact and Befouled Nodes

- **Definition:** A node  $v$  in an FBAS is **intact** iff there exists a DSet  $B$  containing all ill-behaved nodes such that  $v \notin B$
- An optimal FBAS should guarantee safety/liveness for every intact node
- **Definition:** A node  $v$  in an FBAS is **befouled** iff it is not intact
- **Theorem:** In an FBAS with quorum intersection, the set of befouled nodes is a DSet
  - Proof follows from a theorem which says that intersection of DSets is a DSet in an FBAS with quorum intersection



## Federated Voting

# Voting and Ratification

- **Definition:** A node  $v$  **votes** for a statement  $A$  if and only if
  1.  $v$  asserts  $A$  is valid and consistent with all statements  $v$  has accepted, and
  2.  $v$  asserts that it has never voted against  $A$  and promises to not vote against  $A$  in the future.
- **Definition:** A quorum  $U_A$  **ratifies** a statement  $A$  if and only if every member of  $U_A$  votes for  $A$ . A node  $v$  **ratifies**  $A$  iff  $v$  is a member of a quorum  $U_A$  that ratifies  $A$ .
- Theorems
  - Two contradictory statements  $A$  and  $\bar{A}$  cannot both be ratified in an FBAS that enjoys quorum intersection and contains no ill-behaved nodes.
  - Let  $\langle \mathbf{V}, \mathbf{Q} \rangle$  be an FBAS enjoying quorum intersection despite  $B$  where  $B$  contains all ill-behaved nodes. Let  $v_1, v_2 \notin B$ . If  $v_1$  ratifies  $A$ , then  $v_2$  cannot ratify  $\bar{A}$ .
  - Two intact nodes in an FBAS with quorum intersection cannot ratify contradictory statements.

# Accepting Statements

- **Definition:** Let  $v \in \mathbf{V}$  be a node in FBAS  $\langle \mathbf{V}, \mathbf{Q} \rangle$ . A set  $B \subseteq \mathbf{V}$  is  **$v$ -blocking** iff it overlaps with every one of  $v$ 's slices
- **Theorem:** Let  $B \subseteq \mathbf{V}$  be a set of nodes in FBAS  $\langle \mathbf{V}, \mathbf{Q} \rangle$ .  $\langle \mathbf{V}, \mathbf{Q} \rangle$  enjoys quorum availability despite  $B$  iff  $B$  is not  $v$ -blocking for any  $v \in \mathbf{V} \setminus B$ .
- **Corollary:** The DSet of befouled nodes is not  $v$ -blocking for any intact  $v$ .
- **Definition:** An FBAS node  $v$  **accepts** a statement  $A$  iff it has never accepted a statement contradicting  $A$  and it determines that either
  1. There exists a quorum  $U$  such that  $v \in U$  and each each member of  $U$  either voted for  $A$  or claims to accept  $A$ , **or**
  2. each member of a  $v$ -blocking set claims to accept  $A$ .
- Second condition allows  $v$  to vote for  $A$  but later accept  $\bar{A}$
- **Theorem:** Two intact nodes in an FBAS that enjoys quorum intersection cannot accept contradictory statements.

# Confirming Statements

- **Definition:** A quorum  $U_A$  in an FBAS **confirms** a statement  $A$  if and only if every member of  $U_A$  claims to accept  $A$ . A node  $v$  **confirms**  $A$  if and only if it is in such a quorum.
- **Theorem:** Let  $\langle \mathbf{V}, \mathbf{Q} \rangle$  be an FBAS enjoying quorum intersection despite  $B$  where  $B$  contains all ill-behaved nodes. Let  $v_1, v_2 \notin B$ . If  $v_1$  confirms  $A$ , then  $v_2$  cannot confirm  $\bar{A}$ .
- **Theorem:** If an intact node in an FBAS  $\langle \mathbf{V}, \mathbf{Q} \rangle$  with quorum intersection confirms a statement  $A$ , then, whatever subsequently transpires, once sufficient messages are delivered and processed, every intact node will accept and confirm  $A$ .
- But the protocol may get stuck before an intact node confirmation
- Need multiple rounds for liveness

# Stellar Consensus Protocol

- Two subprotocols
  - Nomination protocol
  - Ballot protocol
- Nodes nominate candidate values for a slot which will converge on a composite value
  - Composite value = Union of transaction sets proposed
- Ballot protocol uses federated voting to commit and abort ballots of composite values

# References

- SCP talk <https://www.youtube.com/watch?v=vmwnhZmEZjc>
- SCP white paper <https://www.stellar.org/papers/stellar-consensus-protocol.pdf>
- *Paxos Made Simple*, Leslie Lamport, <https://lamport.azurewebsites.net/pubs/paxos-simple.pdf>
- *How to Build a Highly Available System Using Consensus*, B. W. Lampson, [https://doi.org/10.1007/3-540-61769-8\\_1](https://doi.org/10.1007/3-540-61769-8_1)
- PBFT paper <http://www.pmg.csail.mit.edu/papers/osdi99.pdf>